

Code, partly anonymized

```
File Edit View Run Kernel Git Tabs Settings Help
+ + + + +
/
Name Last Modified
spark-warehouse a month ago
test 19 minutes ago
tuts 21 days ago
FFHS BD Runtime Analytics.pynb 2 minutes ago

[129]: Import pys
import os

(pys.executable) # pip install matplotlib user
(pys.executable) # pip install pandas user
(pys.executable) # pip install pandas-bokoh user
(pys.executable) # pip install seaborn user

Looking in indexes:
Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/site-packages (3.3.4)
Requirement already satisfied: pillow<2.0 in /usr/local/lib/python3.6/site-packages (from matplotlib) (8.2.0)
Requirement already satisfied: cycler<0.10 in /usr/local/lib/python3.6/site-packages (from matplotlib) (0.10.0)
Requirement already satisfied: python-dateutil<2.8 in /usr/local/lib/python3.6/site-packages (from matplotlib) (2.8.1)
Requirement already satisfied: numpy<1.15 in /usr/local/lib/python3.6/site-packages (from matplotlib) (1.19.5)
Requirement already satisfied: pyparsing<2.0.3 in /usr/local/lib/python3.6/site-packages (from matplotlib) (2.1.1)
Requirement already satisfied: six in /usr/local/lib/python3.6/site-packages (from cycler>=0.10.0-matplotlib) (1.16.0)

Looking in indexes:
Requirement already satisfied: pandas in /usr/local/lib/python3.6/site-packages (1.1.1)
Requirement already satisfied: numpy>=1.4 in /usr/local/lib/python3.6/site-packages (from pandas) (2.8.1)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/site-packages (from pandas) (2021.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/site-packages (from python-dateutil>=2.7.3-pandas) (1.16.0)

Looking in indexes:
Requirement already satisfied: pandas-bokoh in /usr/local/lib/python3.6/site-packages (0.2.2)
Requirement already satisfied: pandas-bokoh>=2.0 in /usr/local/lib/python3.6/site-packages (from pandas-bokoh) (1.1.5)
Requirement already satisfied: bokoh>=2.0 in /usr/local/lib/python3.6/site-packages (from pandas-bokoh) (2.3.2)
Requirement already satisfied: pillow>=8.0 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (8.2.0)
Requirement already satisfied: numpy>=1.13 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (1.19.5)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (2.8.1)
Requirement already satisfied: packaging>=16.8 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (20.9)
Requirement already satisfied: typing-extensions>=3.7.4 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (3.10.0.0)
Requirement already satisfied: matplotlib>=2.0 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (2.0.1)
Requirement already satisfied: tornado>=5.1 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (6.1)
Requirement already satisfied: pyzmq>=3.5 in /usr/local/lib/python3.6/site-packages (from bokoh>=2.0-pandas-bokoh) (5.4.1)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.6/site-packages (from packaging>=16.8-bokoh>=2.0-pandas-bokoh) (2.0.1)
Requirement already satisfied: pyparsing>=2.6.2 in /usr/local/lib/python3.6/site-packages (from packaging>=16.8-bokoh>=2.0-pandas-bokoh) (2.4.7)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/site-packages (from pandas>=1.0-2.8-pandas-bokoh) (2021.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/site-packages (from python-dateutil>=2.1-bokoh>=2.0-pandas-bokoh) (1.16.0)

Looking in indexes:
Collecting seaborn
Downloading https://seaborn-0.11.1-py3-none-any.whl (285 kB)

Installing collected packages:
Requirement already satisfied: pandas>=0.21 in /usr/local/lib/python3.6/site-packages (from seaborn) (1.1.5)
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.6/site-packages (from seaborn) (1.19.5)
Requirement already satisfied: matplotliblib>=2.2 in /usr/local/lib/python3.6/site-packages (from seaborn) (3.3.4)

Collecting scipy
Downloading https://scipy-1.5.4-cp36-cp36m-manylinux1_x86_64.whl (25.9 MB)

Requirement already satisfied: setuptools in /usr/local/lib/python3.6/site-packages (from matplotliblib>=2.2-seaborn) (44.1.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.6/site-packages (from matplotliblib>=2.2-seaborn) (0.10.0)
Requirement already satisfied: pyparsing>=2.0.4 in /usr/local/lib/python3.6/site-packages (from matplotliblib>=2.2-seaborn) (2.4.7)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6/site-packages (from matplotliblib>=2.2-seaborn) (2.8.1)
Requirement already satisfied: six in /usr/local/lib/python3.6/site-packages (from python-dateutil>=2.1-seaborn) (1.16.0)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/site-packages (from pandas>=0.23-seaborn) (2021.1)

Installing collected packages:
scipy, seaborn
Successfully installed scipy-1.5.4-seaborn-0.11.1
```

```
Terminal 1 X FFHS BD Runtime Analytics
[1]: # 'os._exit(0)' is optional to force the IPython kernel restart. This is only required when installing matplotlib the first time
# alternatively the IPython kernel may be restarted from the Jupyterhub menu
os._exit(0)

[1]: # Use 'inline' backend to include matplotlib graphs in the notebook
%matplotlib inline

import matplotlib
import matplotlib.pyplot as plt
import pandas as pd
import pandas_bokeh
from bokeh.io import output_notebook, show
from bokeh.plotting import figure

# Pandas config for better display
pd.set_option('display.max_colwidth', 300)

[2]: from pyspark.sql import SparkSession

# Start a PySpark session on the YARN cluster (will take approx. 30 seconds to start)
spark = (SparkSession
        .builder
        .appName("batch-runtime")
        .getOrCreate())

2: Load and clean up data from HDFS

[3]: # data path
hdfs_path = "/user/

Objects data

[4]: obj_file = f"{hdfs_path}/objects.csv"

# use ddl schema
obj_schema = "obj_id INT, obj_name STRING, obj_type_abbr STRING, obj_type_name STRING"

objdf = (spark.read.format("csv")
        .option("header", "true")
        .option("delimiter", ",")
        .schema(obj_schema)
        .load(obj_file))

objdf.printSchema()
print((objdf.count(), len(objdf.columns)))
print(objdf.show())

root
|-- obj_id: integer (nullable = true)
|-- obj_name: string (nullable = true)
|-- obj_type_abbr: string (nullable = true)
|-- obj_type_name: string (nullable = true)

(77683, 4)
+-----+-----+-----+-----+
|obj_id|          obj_name|obj_type_abbr|obj_type_name|
+-----+-----+-----+-----+
| 19497|      SR finisObs 03|          3083|          Job|
| 19498|SR Prepare Manual...|          3083|          Job|
| 19499|SR Update Tagessc...|          3083|          Job|
| 19500|BU Buchungstransa...|          3083|          Job|
| 19502|BU Manueller Abbr...|          3083|          Job|
+-----+-----+-----+-----+

only showing top 5 rows
```

```
Terminal 1  FFHS BD Runtime Analytics.iX

Runtime data

[5]: rt_file = f'{hdfs_path}/runtime.csv'

rtdf = (spark.read.format('csv')
        .option('inferSchema', 'true')
        .option('header', 'true')
        .option('delimiter', ';')
        .load(rt_file))

# Convert columns to lower case (upper case caused by inferSchema option)
for col in rtdf.columns:
    rtdf = rtdf.withColumnRenamed(col, col.lower())

rtdf.printSchema()
print((rtdf.count(), len(rtdf.columns)))

root
|-- obj_id: integer (nullable = true)
|-- obj_emv_id: long (nullable = true)
|-- comp_id: integer (nullable = true)
|-- stats_id: long (nullable = true)
|-- stats_rid: integer (nullable = true)
|-- stats_day: integer (nullable = true)
|-- stats_act: string (nullable = true)
|-- stats_start: string (nullable = true)
|-- stats_end: string (nullable = true)
|-- stats_dur_s: integer (nullable = true)
|-- stats_status: string (nullable = true)
|-- stats_arch_1: string (nullable = true)
|-- stats_arch_2: string (nullable = true)

(9345043, 13)

[6]: from pyspark.sql.functions import *

# convert string timestamp to date
rtdf = rtdf.withColumn('stats_act_date', to_date(col('stats_act'), 'dd.MM.yyyy')).drop('stats_act')
rtdf = rtdf.withColumn('stats_start_date', to_date(col('stats_start'), 'dd.MM.yyyy')).drop('stats_start')
rtdf = rtdf.withColumn('stats_end_date', to_date(col('stats_end'), 'dd.MM.yyyy')).drop('stats_end')

rtdf.printSchema()
rtdf.show(5)

root
|-- obj_id: integer (nullable = true)
|-- obj_emv_id: long (nullable = true)
|-- comp_id: integer (nullable = true)
|-- stats_id: long (nullable = true)
|-- stats_rid: integer (nullable = true)
|-- stats_day: integer (nullable = true)
|-- stats_dur_s: integer (nullable = true)
|-- stats_status: string (nullable = true)
|-- stats_arch_1: string (nullable = true)
|-- stats_arch_2: string (nullable = true)
|-- stats_act_date: date (nullable = true)
|-- stats_start_date: date (nullable = true)
|-- stats_end_date: date (nullable = true)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| obj_id|obj_emv_id|comp_id| stats_id | stats_rid|stats_day|stats_dur_s|stats_status|stats_arch_1|stats_arch_2|stats_act_date|stats_start_date|stats_end_date|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 346280688| 688216355|      6|8151814549|1341181044|      3|      171|  ENDED OK|      AV|  AVALOQ_JOB|2019-12-31|  2020-01-01|  2020-01-01|
| 342658188| 679440702|      7|8151839639|1341069383|      3|      352|  ENDED OK|      EOD|  AVALOQ_JOB|2019-12-31|  2020-01-01|  2020-01-01|
| 376661047| 747477778|      7|8151844866|1341052877|      3|      3939|  ENDED OK|      EOD|  AVALOQ_JOB|2019-12-31|  2020-01-01|  2020-01-01|
| 9198662| 2267939|      9|8151852272|1341113251|      3|      81|  ENDED OK|      EOD|  AVALOQ_JOB|2019-12-31|  2020-01-01|  2020-01-01|
| 9195308| 2265140|      5|8151857252|1341184712|      3|      444|  ENDED OK|      EOD|  AVALOQ_JOB|2019-12-31|  2020-01-01|  2020-01-01|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
Terminal 1  FFHS BD Runtime Analytics.iX

Planning data

[7]: plan_file = f'{hdfs_path}/planning.csv'

# use ddl schema
plan_schema = 'plan_key STRING, plan_id INT, plan_status STRING, plan_type STRING, plan_subtype_1 STRING, plan_start STRING, plan_end STRING'

plandf = (spark.read.format('csv')
        .option('inferSchema', 'true')
        .option('header', 'true')
        .option('delimiter', ';')
        .option('encoding', 'UTF-8')
        .schema(plan_schema)
        .load(plan_file))

plandf.printSchema()
print((plandf.count(), len(plandf.columns)))

root
|-- plan_key: string (nullable = true)
|-- plan_id: integer (nullable = true)
|-- plan_status: string (nullable = true)
|-- plan_type: string (nullable = true)
|-- plan_subtype_1: string (nullable = true)
|-- plan_start: string (nullable = true)
|-- plan_end: string (nullable = true)

(47, 7)

[8]: from pyspark.sql.functions import *

# convert string timestamp to date
plandf = plandf.withColumn('plan_start_date', to_date(col('plan_start'), 'dd.MM.yyyy')).drop('plan_start')
plandf = plandf.withColumn('plan_end_date', to_date(col('plan_end'), 'dd.MM.yyyy')).drop('plan_end')

plandf.printSchema()
plandf.show(5)

root
|-- plan_key: string (nullable = true)
|-- plan_id: integer (nullable = true)
|-- plan_status: string (nullable = true)
|-- plan_type: string (nullable = true)
|-- plan_subtype_1: string (nullable = true)
|-- plan_start_date: date (nullable = true)
|-- plan_end_date: date (nullable = true)

+-----+-----+-----+-----+-----+-----+
| plan_key|plan_id|plan_status|      plan_type|plan_subtype_1|plan_start_date|plan_end_date|
+-----+-----+-----+-----+-----+-----+
| PLAN-5336|3825484|  Resolved|      Release|      Core|  2021-06-05|  2021-06-06|
| PLAN-4675|2570704|  Resolved|  Maintenance|      CPW|  2021-05-15|  2021-05-16|
| PLAN-4378|2683194|  Resolved|Maintenance Emerg...|      CPW|  2021-05-07|  2021-05-07|
| PLAN-5413|3868011|  Resolved|      Supply|      Core|  2021-04-29|  2021-04-29|
| PLAN-5281|2993294|  Resolved|      Supply|      Core|  2021-04-15|  2021-04-15|
+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

3: End-of-Period Master Analysis

Data Preparation

```
[9]: # join objects and runtime data
# filter End-of-Period workflow
# filter one customer only, customer id narrowed for anonymization purpose
# workdays only, ignore weekends

eopdf = rtdf \
    .join(objdf, objdf.obj_id == rtdf.obj_id) \
    .filter(objdf.obj_name == 'End-of-Period workflow') \
    .filter(rtdf.comp_id == 2) \
    .filter(rtdf.stats_day < 6) \
    .select(rtdf.stats_start_date, rtdf.stats_dur_s) \
    .sort(rtdf.stats_start_date)

print(eopdf.count())
eopdf.show(25)
```

```
367
+-----+
|stats_start_date|stats_dur_s|
+-----+
| 2020-01-01 | 12447 |
| 2020-01-02 | 44711 |
| 2020-01-03 | 45179 |
| 2020-01-06 | 45678 |
| 2020-01-07 | 44015 |
| 2020-01-08 | 47014 |
| 2020-01-09 | 43947 |
| 2020-01-10 | 44532 |
| 2020-01-13 | 44694 |
| 2020-01-14 | 44631 |
| 2020-01-15 | 44353 |
| 2020-01-16 | 44409 |
| 2020-01-17 | 44480 |
| 2020-01-20 | 44657 |
| 2020-01-21 | 44214 |
| 2020-01-22 | 44425 |
| 2020-01-23 | 45984 |
| 2020-01-24 | 45937 |
| 2020-01-27 | 46150 |
| 2020-01-28 | 45383 |
| 2020-01-29 | 45348 |
| 2020-01-30 | 46522 |
| 2020-01-31 | 54226 |
| 2020-02-03 | 45018 |
| 2020-02-04 | 44881 |
+-----+
only showing top 25 rows
```

```
[10]: from pyspark.sql.functions import lit, col

# get max runtime duration and add it as new static column to the planning dataframe. will be used for plotting later
max_dur = eopdf.agg(lit("stats_dur_s": "max")).collect()
print(max_dur)

# to-do: get value from variable
# plandf = plandf.withColumn('yhelper', lit(col(max_dur)))
plandf = plandf.withColumn('yhelper', lit(73000))

plandf.show(5)

[Row(max(stats_dur_s)=73000)]
+-----+
|plan_key|plan_id|plan_status|plan_type|plan_subtype_1|plan_start_date|plan_end_date|yhelper|
+-----+
|PLAN-5336|3825484|Resolved|Release|Core|2021-06-05|2021-06-06|73000|
|PLAN-4675|2570744|Resolved|Maintenance|CPM|2021-05-15|2021-05-16|73000|
|PLAN-4378|2683194|Resolved|Maintenance Emerg...|CPM|2021-05-07|2021-05-07|73000|
|PLAN-5413|3860011|Resolved|Supply|Core|2021-04-29|2021-04-29|73000|
|PLAN-5281|2993294|Resolved|Supply|Core|2021-04-15|2021-04-15|73000|
+-----+
only showing top 5 rows
```

```
[11]: # convert to Pandas (for easy plotting)
eopdf_pd = eopdf.toPandas()
plandf_pd = plandf.toPandas()
```


Plotting

```
[12]: plt.figure(figsize=(12, 6), dpi=80)

plt.plot(eopdf_pd.stats_start_date, eopdf_pd.stats_dur_s, label='label 1')
plt.plot(plandf_pd.plan_end_date, plandf_pd.yhelper, 'bo', label='label 2')

plt.title('EOP Master Analysis')
plt.ylabel('Runtime')
plt.xlabel('Date')

plt.show()
```



```
Terminal 1  FFHS BD Runtime Analytics
Quick Analysis:
looks like runtime increased between approx. mid of april - mid may

[13]: import pandas as pd
      from pyspark.sql.functions import *

      # get avg runtime duration for comparison
      avg_dur = eopdf.agg({"stats_dur_s": "avg"}).collect()
      print(avg_dur)

      # set filter for affected time period
      plan_df.filter(col("plan_end_date").between(pd.to_datetime("2021-03-01"),pd.to_datetime("2021-05-30"))).show()
      eopdf.filter(col("stats_start_date").between(pd.to_datetime("2021-04-01"),pd.to_datetime("2021-05-15"))).show(100)

      [Row(avg(stats_dur_s)=47199.42779291553)]

      +-----+-----+-----+-----+-----+-----+-----+-----+
      | plan_key|plan_id|plan_status|      plan_type|plan_subtype|plan_start_date|plan_end_date|yhelper|
      +-----+-----+-----+-----+-----+-----+-----+-----+
      |PLAN-4675|2570744|Resolved|Maintenance|CPMM|2021-05-15|2021-05-16|73000|
      |PLAN-4378|2683194|Resolved|Maintenance Emerg...|CPMM|2021-05-07|2021-05-07|73000|
      |PLAN-5413|3060011|Resolved|Supply|Core|2021-04-29|2021-04-29|73000|
      |PLAN-5281|2993294|Resolved|Supply|Core|2021-04-15|2021-04-15|73000|
      |PLAN-4674|2570743|Resolved|Maintenance|GPMW|2021-04-10|2021-04-11|73000|
      |PLAN-4900|2626684|Resolved|Supply|Core|2021-03-18|2021-03-18|73000|
      |PLAN-4673|2570741|Resolved|Maintenance|GPMW|2021-03-13|2021-03-14|73000|
      |PLAN-5277|2993266|Resolved|Supply|Core|2021-03-06|2021-03-06|73000|
      |PLAN-5337|3026056|Resolved|Release|Core|2021-03-06|2021-03-06|73000|
      |PLAN-5280|2993290|Resolved|Supply|Core|2021-03-04|2021-03-04|73000|
      +-----+-----+-----+-----+-----+-----+-----+-----+

      +-----+-----+
      |stats_start_date|stats_dur_s|
      +-----+-----+
      |2021-04-02|45436|
      |2021-04-05|43703|
      |2021-04-06|48047|
      |2021-04-07|49839|
      |2021-04-08|47350|
      |2021-04-09|46956|
      |2021-04-12|50515|
      |2021-04-13|50836|
      |2021-04-14|50350|
      |2021-04-15|50885|
      |2021-04-16|49796|
      |2021-04-19|49041|
      |2021-04-20|49600|
      |2021-04-21|49131|
      |2021-04-22|53039|
      |2021-04-23|53733|
      |2021-04-26|50387|
      |2021-04-27|50725|
      |2021-04-28|51246|
      |2021-04-29|53610|
      |2021-04-30|66974|
      |2021-05-03|57213|
      |2021-05-04|53721|
      |2021-05-05|51415|
      |2021-05-06|58195|
      |2021-05-07|47167|
      |2021-05-10|47485|
      |2021-05-11|47683|
      |2021-05-12|47150|
      |2021-05-13|48886|
      |2021-05-14|47004|
      +-----+-----+
```

```
Terminal 1  FFHS BD Runtime Analytics
4: Slow Poison Analysis

[14]: from pyspark.sql.functions import *

      # prepared dataframes from above: objdf, rtidf, plandf

      # define static parameters
      min_rec = 3 # minimum number of records
      rt_status = "ENDED_OK" # final status, ignore aborted runs
      quarters = 4 # number of time periods to check

      # TEMP for testing only -> also remove in filters below
      nl = [9193523,9197444,9197403]
      # filter(rtidf_cnt.obj_id.isin(nl)) \
      # TEMP end

      # set timeframe
      # to-do: calculate dates based on current date and defined parameters
      date_start = "2020-07-01"
      date_end = "2021-06-30"

      # count no. of records per ID, rename "count"
      rtidf_cnt = rtidf.join(rtidf.groupby("obj_id").count().withColumnRenamed("count","runs"),on="obj_id")

      print(rtidf_cnt.count())
      rtidf_cnt.show(5)

      # set date and customer filters, objects >= the minimum no. of records. necessary to compute median later
      rtidf_base = rtidf_cnt \
        .filter(rtidf_cnt.stats_act_date.between(date_start,date_end)) \
        .filter(rtidf_cnt.comp_id == 0) \
        .filter(rtidf_cnt.stats_status == rt_status) \
        .filter(rtidf_cnt.runs >= min_rec) \
        .drop("obj_omv_id", "comp_id", "stats_id", "stats_rid", "stats_day", "comp_id", "stats_status", "stats_arch_1", "stats_arch_2", "stats_start_date", "stats_end_date", "runs")

      print(rtidf_base.count())
      rtidf_base.show(5)

      9345043

      +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
      | obj_id|obj_omv_id|comp_id| stats_id|stats_rid|stats_day|stats_dur_s|stats_status|stats_arch_1|stats_arch_2|stats_act_date|stats_start_date|stats_end_date|runs|
      +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
      |9170264|2252670|7|12426787162|1815059007|4|352|ENDED_OK|EOD|AWALQ_300|2021-06-02|2021-06-03|2021-06-03|199|
      |9170264|2252670|7|12440520087|1821030440|4|768|ENDED_OK|EOD|AWALQ_300|2021-06-09|2021-06-10|2021-06-10|199|
      |9170264|2252670|7|12451678675|1823078088|6|277|ENDED_OK|EOD|AWALQ_300|2021-06-11|2021-06-12|2021-06-12|199|
      |9170264|2252670|7|12455548453|1822057274|5|450|ENDED_OK|EOD|AWALQ_300|2021-06-10|2021-06-11|2021-06-11|199|
      |9170264|2252670|7|12286237270|1790992464|4|68|ENDED_OK|EOD|AWALQ_300|2021-05-05|2021-05-06|2021-05-06|199|
      +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

      only showing top 5 rows

      303100

      +-----+-----+-----+
      | obj_id|stats_dur_s|stats_act_date|
      +-----+-----+-----+
      |9195811|80|2020-10-02|
      |9195811|78|2020-10-05|
      |9195811|93|2021-02-12|
      |9195811|66|2020-07-20|
      |9195811|69|2020-07-30|
      +-----+-----+-----+

      only showing top 5 rows
```

```
Terminal 1 X FFHS BD Runtime Analytics X
[15]: from pyspark.sql.types import IntegerType
from pyspark.sql import functions as F
from pyspark.sql import SparkSession
from pyspark.sql.types import IntegerType
import numpy as np

#TEMP
rtdf_base_backup = rtdf_base
# rtdf_base = rtdf_base_backup
#TEMP end

rtdf_base = rtdf_base.withColumn('quarters',concat(F.lit('_'),F.year(rtdf_base.stats_act_date),F.lit('_'),F.quarter(rtdf_base.stats_act_date)).cast('string')).drop('stats_act_date')
rtdf_base.show(15)

# calculate median for defined time periods and create pivot table
# original code -> does not work because percentile_approx requires Spark >= 3.1 (current version installed 3.0)
# output = rtdf_base.groupby('id').pivot('quarters').agg(
#     percentile_approx('stats_dur_s', 0.5, F.lit(100000)))

# workaround: convert to Pandas
df_pd = rtdf_base.toPandas()
#print(df_pd)
#display(df_pd)

# create pivot table with median for each time period
# convert to integer (for whatever reason, the last median value had data type LONG)
df_pd = pd.pivot_table(df_pd, values='stats_dur_s', fill_value=0, index=['obj_id'], columns=['quarters'], aggfunc=np.median).astype(int)
#print(df_pd_out)
#display(df_pd)

# convert back to Spark
rtdf_piv = spark.createDataFrame(df_pd.reset_index(drop=False))
# rtdf_piv.show()
# rtdf_piv.printSchema()

# calculate difference between time periods (seconds and percent)
# to-do: dynamic columns -> use column id instead of fixed names
# rtdf_piv = rtdf_piv.withColumn('q2_q1_s',rtdf_piv.select(rtdf_piv.columns[2]) - rtdf_piv.select(rtdf_piv.columns[1]))
# rtdf_piv.show(2)
rtdf_piv = rtdf_piv \
    .withColumn('q2_q1_s',F.col('m_2020_3') - F.col('m_2020_1')) \
    .withColumn('q2_q1_p',F.round((F.col('m_2020_3') - F.col('m_2020_1'))/F.col('m_2020_3')*100,1)) \
    .withColumn('q3_q2_s',F.round((F.col('m_2021_1') - F.col('m_2020_4')))) \
    .withColumn('q3_q2_p',F.round((F.col('m_2021_1') - F.col('m_2020_4'))/F.col('m_2020_4')*100,1)) \
    .withColumn('q4_q3_s',F.col('m_2021_2') - F.col('m_2021_1')) \
    .withColumn('q4_q3_p',F.round((F.col('m_2021_2') - F.col('m_2021_1'))/F.col('m_2021_1')*100,1))

rtdf_piv.show(15)

-----+-----+
| obj_id|stats_dur_s|quarters|
-----+-----+
[9195811]      63|m_2021_2|
[9195811]      78|m_2021_2|
[9195811]      63|m_2021_2|
[9195811]      66|m_2021_2|
[9195811]      80|m_2021_2|
[9195811]      69|m_2021_2|
[9195811]      69|m_2021_2|
[9195811]      62|m_2021_2|
[9195811]      80|m_2021_2|
[9195811]      86|m_2021_2|
[9195811]      77|m_2020_4|
[9195811]      63|m_2020_4|
[9195811]      71|m_2020_4|
[9195811]      68|m_2020_4|
[9195811]      71|m_2020_4|
-----+-----+
only showing top 15 rows
```

```
Terminal 1 X FFHS BD Runtime Analytics X
-----+-----+
| obj_id|m_2020_3|m_2020_4|m_2021_1|m_2021_2|q2_q1_s|q2_q1_p|q3_q2_s|q3_q2_p|q4_q3_s|q4_q3_p|
-----+-----+
[9163443]    125|    87|    0|   112|   -38|  -30.4|   -87|  -100.0|   112|   null|
[9165254]    150|    0|  10121|    0|  -150| -100.0|  10121|   null| -10121| -100.0|
[9165709]    278|    0|    87|   114|  -278| -100.0|    87|   null|    27|  31.0|
[9165790]    77|    0|    0|    0|   -77| -100.0|    0|   null|    0|   null|
[9165930]   167|  100|   106|   114|  -79|  -42.2|   -2|   -1.9|    8|   7.5|
[9166000]    75|    0|    71|    61|  -75| -100.0|   71|   null|   -10| -14.1|
[9166073]    75|    0|    65|    0|  -79| -100.0|   65|   null|   -65| -100.0|
[9166156]    85|   92|    0|    73|    7|    8.2|  -92| -100.0|    73|   null|
[9166930]   123|  151|   149|   327|   28|   22.8|   -2|   -1.3|   178|  119.5|
[9167229]   111|  105|   102|   106|   -6|   -5.4|   -3|   -2.9|    4|   3.9|
[9167233]   111|  106|   100|   105|   -5|   -4.5|   -6|   -5.7|    5|   5.0|
[9171199]  49012| 50315| 49836| 51457| 1303|    2.7|  -479|   -1.0| 1621|    3.3|
[9171275]   9831| 10254| 10731| 10805|  423|    4.3|   477|    4.7|    74|    0.7|
[9171306]   117|   110|   117|   118|   -7|   -6.0|    7|    6.4|    1|    0.9|
[9171319]   2432|  2431|  2428|  2429|   -1|    0.0|   -3|   -0.1|    1|    0.0|
-----+-----+
only showing top 15 rows

[16]: # define runtime parameters
min_rt = 600 # minimum runtime in seconds
min_rt_incr = 3 # minimum increase of runtime in percent

# filter relevant runtime records
rtdf_out = rtdf_piv \
    .filter(rtdf_piv.m_2021_2 >= min_rt) \
    .filter((rtdf_piv.q2_q1_p >= min_rt_incr) & (rtdf_piv.q3_q2_p >= min_rt_incr) & (rtdf_piv.q4_q3_p >= min_rt_incr)) \
    .select(rtdf_piv.obj_id, rtdf_piv.m_2020_3, rtdf_piv.m_2020_4, rtdf_piv.m_2021_1, rtdf_piv.m_2021_2)

rtdf_out.show(15)

-----+-----+
| obj_id|m_2020_3|m_2020_4|m_2021_1|m_2021_2|
-----+-----+
[9190370]    3070|   3392|   3707|   3951|
[9193332]    1559|   1690|   1786|   1962|
[9193504]    1447|   1703|  12224|  12952|
[9193523]    1446|   1701|  12220|  12951|
[9195801]    467|   604|   636|   667|
[9195975]    855|   923|   1013|   1241|
[9196060]   2873|  3310|  3411|  3532|
[9197403]    609|   645|   771|   795|
[9197404]   1176|  1304|  1356|  1403|
[9197594]    309|   436|   903|   938|
[9197642]   2313|  2411|  2579|  2904|
[9197643]    984|  1131|  1314|  1409|
[9198715]   8378|  10708|  11470|  11994|
[9198719]   8377|  10707|  11470|  11976|
[9199099]    274|   339|   605|   656|
-----+-----+
only showing top 15 rows

[17]: rtdf_out.printSchema()

root
 |-- obj_id: long (nullable = true)
 |-- m_2020_3: long (nullable = true)
 |-- m_2020_4: long (nullable = true)
 |-- m_2021_1: long (nullable = true)
 |-- m_2021_2: long (nullable = true)
```

```
Terminal 1  FFHS BD Runtime Analytics.i
|-- m_2021_2: long (nullable = true)

[18]: from pyspark.sql.functions import expr

unpivotExpr = "stack(4, 'm_2020_3', m_2020_3, 'm_2020_4', m_2020_4, 'm_2021_1', m_2021_1, 'm_2021_2', m_2021_2) as (period,med)"
rtdf_up = rtdf_out.select("obj_id", expr(unpivotExpr))

rtdf_up.show()

+-----+-----+-----+
| obj_id | period | med |
+-----+-----+-----+
| 9190370 | m_2020_3 | 3070 |
| 9190370 | m_2020_4 | 3392 |
| 9190370 | m_2021_1 | 3707 |
| 9190370 | m_2021_2 | 3951 |
| 9193332 | m_2020_3 | 1559 |
| 9193332 | m_2020_4 | 1690 |
| 9193332 | m_2021_1 | 1786 |
| 9193332 | m_2021_2 | 1962 |
| 9193504 | m_2020_3 | 1447 |
| 9193504 | m_2020_4 | 1703 |
| 9193504 | m_2021_1 | 12224 |
| 9193504 | m_2021_2 | 12952 |
| 9193523 | m_2020_3 | 1446 |
| 9193523 | m_2020_4 | 1701 |
| 9193523 | m_2021_1 | 12220 |
| 9193523 | m_2021_2 | 12951 |
| 9195801 | m_2020_3 | 467 |
| 9195801 | m_2020_4 | 604 |
| 9195801 | m_2021_1 | 636 |
| 9195801 | m_2021_2 | 667 |
+-----+-----+-----+
only showing top 20 rows

[19]: from pyspark.sql.functions import *
from pyspark.sql import functions as F
from pyspark.sql.functions import when
from pyspark.sql.functions import concat_ws

# set end of quarter dates (mm-dd)
q1_end = '03-31'
q2_end = '06-30'
q3_end = '09-30'
q4_end = '12-31'

# extract year and quarter, creat new date column (last day of time period)
rtdf_final = rtdf_up \
    .withColumn('m_year', substring('period', 3,4)) \
    .withColumn('m_quarter', substring('period', 8,8))

rtdf_final = rtdf_final.withColumn('m_date', when(rtdf_final.m_quarter.endswith('1'), q1_end) \
    .when(rtdf_final.m_quarter.endswith('2'), q2_end) \
    .when(rtdf_final.m_quarter.endswith('3'), q3_end) \
    .when(rtdf_final.m_quarter.endswith('4'), q4_end) \
    .otherwise(rtdf_final.m_quarter))

rtdf_final = rtdf_final \
    .withColumn('end_date', to_date(concat_ws('-', rtdf_final.m_year, rtdf_final.m_date))) \
    .drop('period', 'm_year', 'm_quarter', 'm_date')

# join with objects data
rtdf_final = rtdf_final \
    .join(objdf, objdf.obj_id == rtdf_final.obj_id) \
    .filter(objdf.obj_type_abbr == "JOBS") \
    .select(objdf.obj_name, rtdf_final.end_date, rtdf_final.med) \
    .sort(objdf.obj_name, rtdf_final.end_date)

rtdf_final.show()
```

```
Terminal 1  FFHS BD Runtime Analytics.ij
rtddf_final.show()
rtddf_final.printSchema()

+-----+-----+-----+
| obj_name | end_date | med |
+-----+-----+-----+
| ACCT.2 | 2020-09-30 | 855 |
| ACCT.2 | 2020-12-31 | 923 |
| ACCT.2 | 2021-03-31 | 1013 |
| ACCT.2 | 2021-06-30 | 1241 |
| AVQ.2 | 2020-09-30 | 3070 |
| AVQ.2 | 2020-12-31 | 3392 |
| AVQ.2 | 2021-03-31 | 3707 |
| AVQ.2 | 2021-06-30 | 3951 |
| COMPL | 2020-09-30 | 637 |
| COMPL | 2020-12-31 | 1482 |
| COMPL | 2021-03-31 | 4372 |
| COMPL | 2021-06-30 | 5359 |
| CRED.2 | 2020-09-30 | 609 |
| CRED.2 | 2020-12-31 | 645 |
| CRED.2 | 2021-03-31 | 771 |
| CRED.2 | 2021-06-30 | 795 |
| CRED.2 | 2020-09-30 | 1176 |
| CRED.2 | 2020-12-31 | 1304 |
| CRED.2 | 2021-03-31 | 1356 |
| CRED.2 | 2021-06-30 | 1403 |
+-----+-----+-----+
only showing top 20 rows

root
|-- obj_name: string (nullable = true)
|-- end_date: date (nullable = true)
|-- med: long (nullable = true)

[20]: # convert to Pandas (for easy plotting)
rtddf_plot_pd = rtddf_final.toPandas()

[21]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 4), dpi=80)

# create plot
sns.set(style="darkgrid")

sns.lineplot(
    data=rtddf_plot_pd,
    x='end_date', y='med', hue='obj_name', style='obj_name', legend='full',
    markers=True, dashes=False)

plt.title('Slow Poison Analysis', fontsize=16)
plt.xlabel('Date')
plt.ylabel('Runtime [Median in sec.]')

plt.xticks(rotation=45)
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
```

